

Class 7: Introduction to Machine Learning Lab

Alisa Zhang (PID: A18299618)

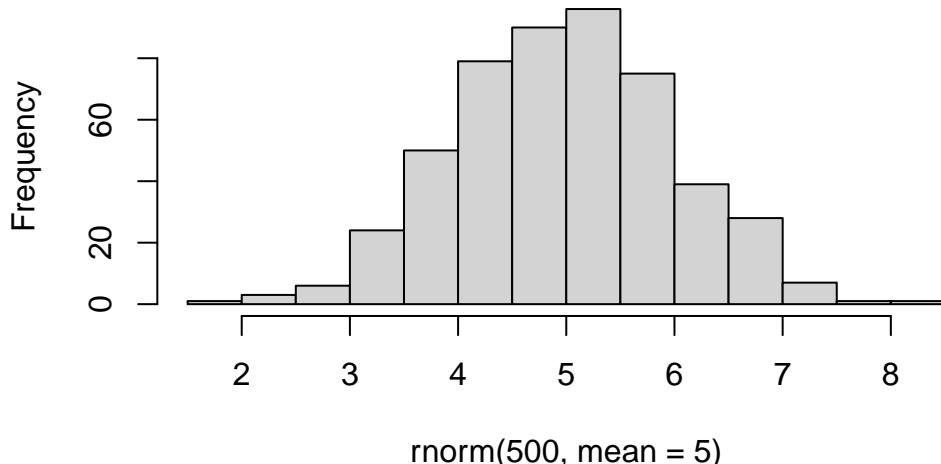
Today we will explore some fundamental machine learning methods including clustering and dimensionality reduction.

K-means clustering

To see how this works, let's first makeup some data to cluster where we know what the answer should be. We can use the `rnorm()` fuction to help here:

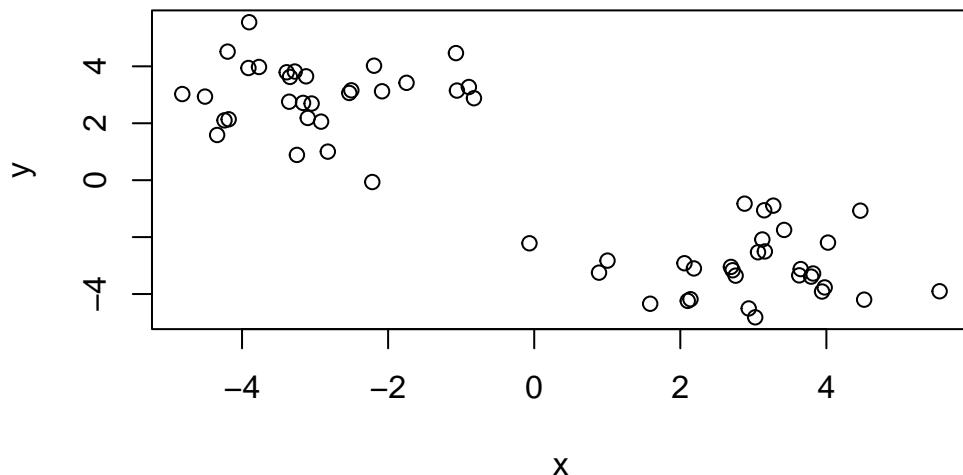
```
hist(rnorm(500, mean=5))
```

Histogram of `rnorm(500, mean = 5)`



```
x <- c(rnorm(30, mean=-3), rnorm(30, mean=3))
y <- rev(x)
```

```
x <- cbind(x, y)  
plot(x)
```



The function for K-means clustering in “base” R is `kmeans()`

```
km <- kmeans(x, centers = 2)  
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1 -2.992952  2.981686
2  2.981686 -2.992952

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 74.84455 74.84455  
(between_SS / total_SS =  87.7 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

To get the results of the `returnrend` list object, we can use the dollar \$ syntax

Q1. How many points are in each cluster?

km\$size

[1] 30 30

Q2. What ‘component’ of your result object details - cluster assignment/membership? - cluster center?

km\$cluster

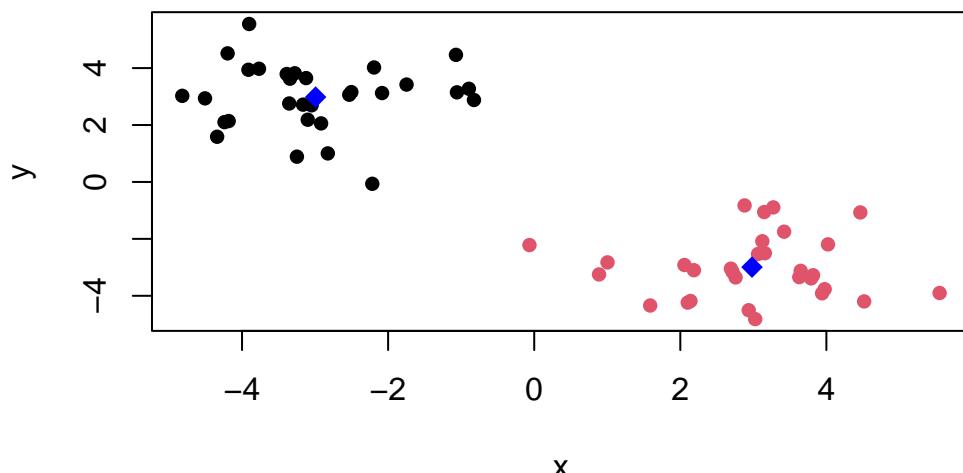
km\$centers

	x	y
1	-2.992952	2.981686
2	2.981686	-2.992952

Q3. Make a clustering result figure of the data colored by cluster membership.

```
plot(x, col=km$cluster, pch=16)
# if col=c("red", "blue"), the vector would recycle to length of 60 alternating colors

points(km$centers, col="blue", pch=18, cex=1.5)
```

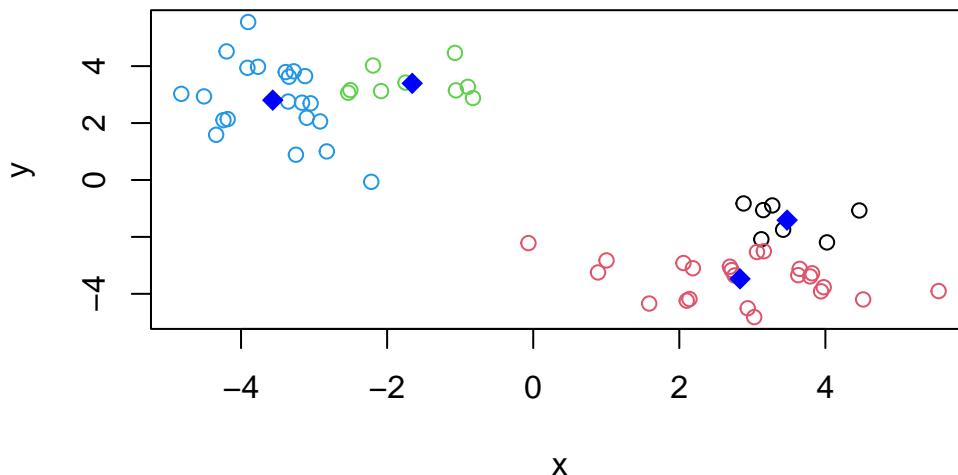


K-means clustering is very popular as it is very fast and relatively straightforward: it takes numeric data as input and returns the cluster membership vector etc.

The “issue” is that we tell `kmeans()` how many clusters we want.

Q4. Run kmeans again and cluster into 4 groups and plot the results like we did above.

```
km2 <- kmeans(x, centers=4)
plot(x, col=km2$cluster)
points(km2$centers, col="blue", pch=18, cex=1.5)
```



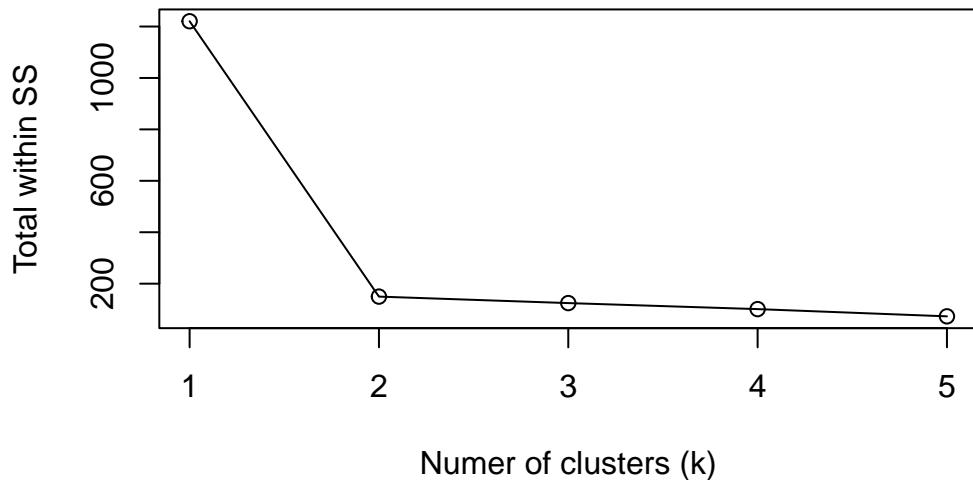
```
# Within Sum of Square
mean(km$withinss)
```

```
[1] 74.84455
```

Scree plot

```
ssvec <- c()
for(i in 1:5){
  km <- kmeans(x, centers = i)
  ssvec[i] <- km$tot.withinss
}
plot(c(1:5), ssvec, type="o", xlab="Numer of clusters (k)", ylab="Total within SS", main="Scree Plot")
```

Scree Plot



Hierarchical Clustering

The main base R function for Hierarchical Clustering is called `hclust()`. Here we can't just input our data, we need to first calculate a distance matrix (e.g. `dist()`) for our data and use this as input to `hclust()`.

```
d <- dist(x)
hc <- hclust(d)
hc
```

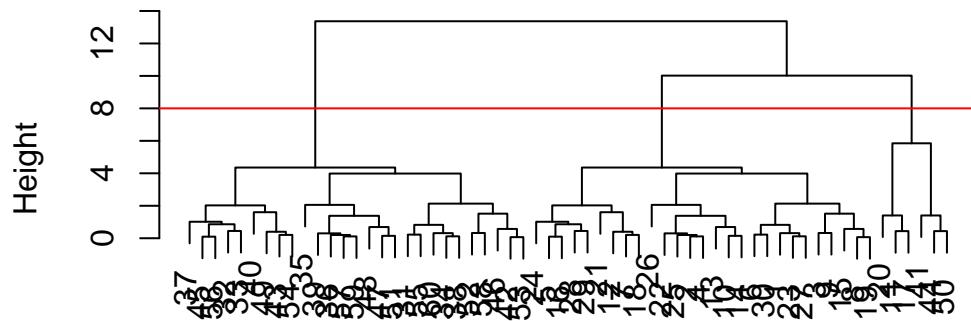
```
Call:
hclust(d = d)

Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

There is a plot method for class `hclust` results. Lets try

```
plot(hc)
abline(h=8, col='red')
```

Cluster Dendrogram



```
d  
hclust (*, "complete")
```

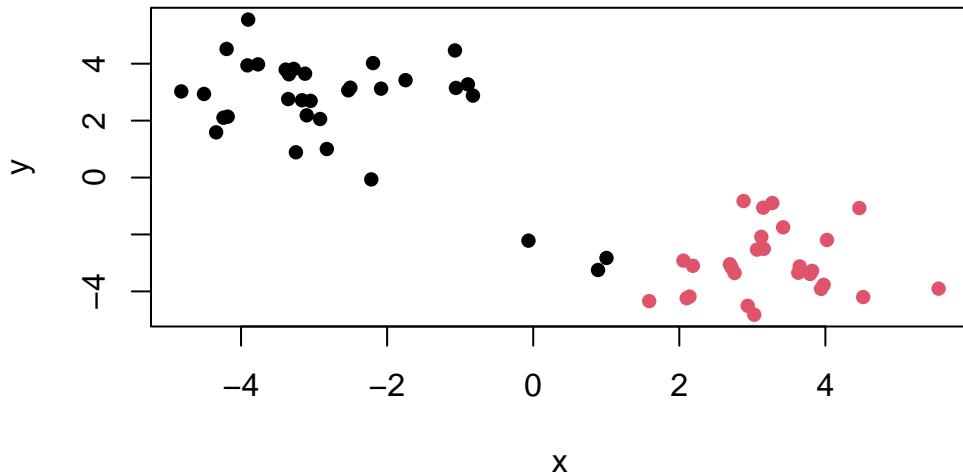
To get our cluster membership vector (i.e. our main cluster result), we can *cut the tree* at a given height or at a height that yields a given “ k ”.

```
cutree(hc, h=8)
```

```
grps <- cutree(hc, k=2)
```

Q. Plot the data with our hclust result coloring

```
plot(x, col = grps, pch = 16)
```



Principal Component Analysis (PCA)

PCA of UK food data

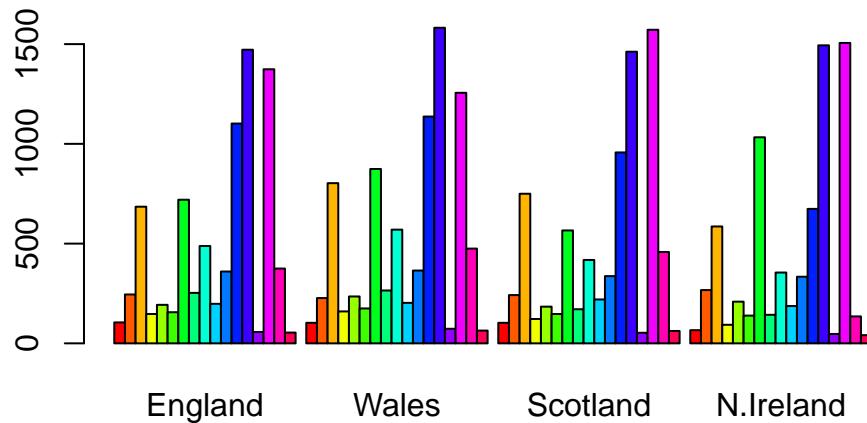
Import food data from an online CSV file:

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

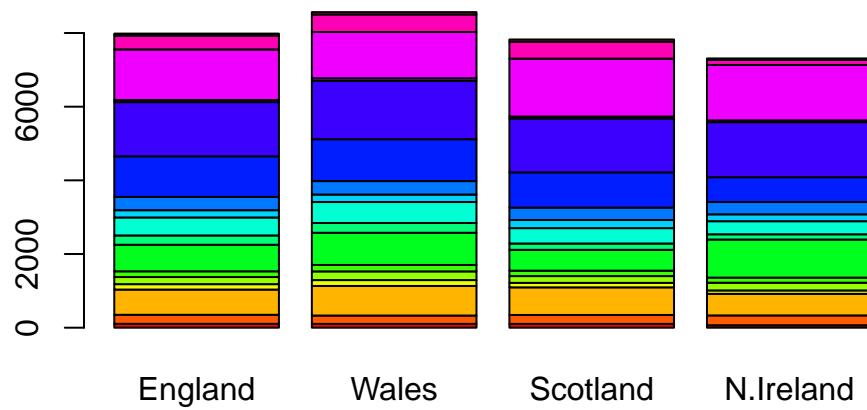
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Some base figures

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

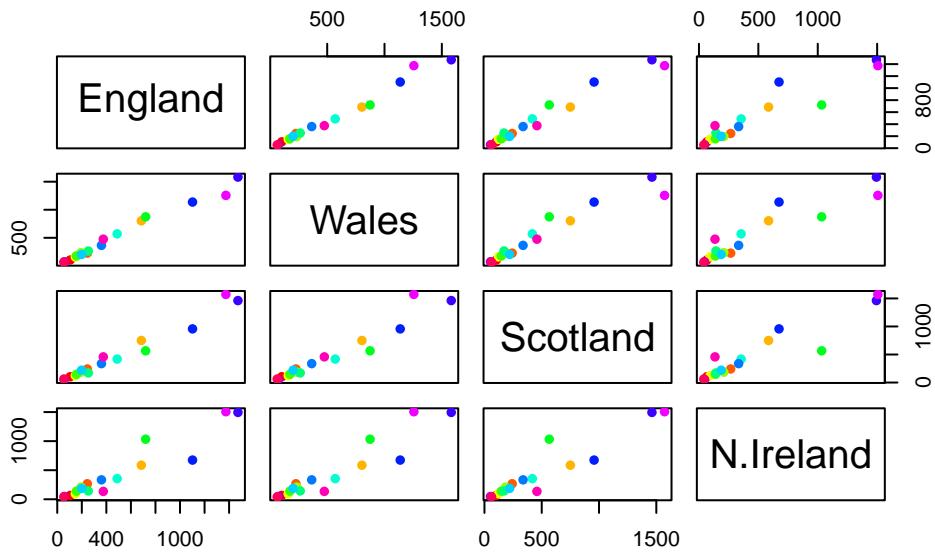


```
# Stacked bar chart - NOT useful
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



There is one plot that can be useful for small datasets:

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



Main point: It can be difficult to spot major trends and patterns even in relatively small multivariate datasets (here we only have 17 dimensions. Typically we have 1000s).

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

I will take the transpose of our data so the “foods” are in columns.

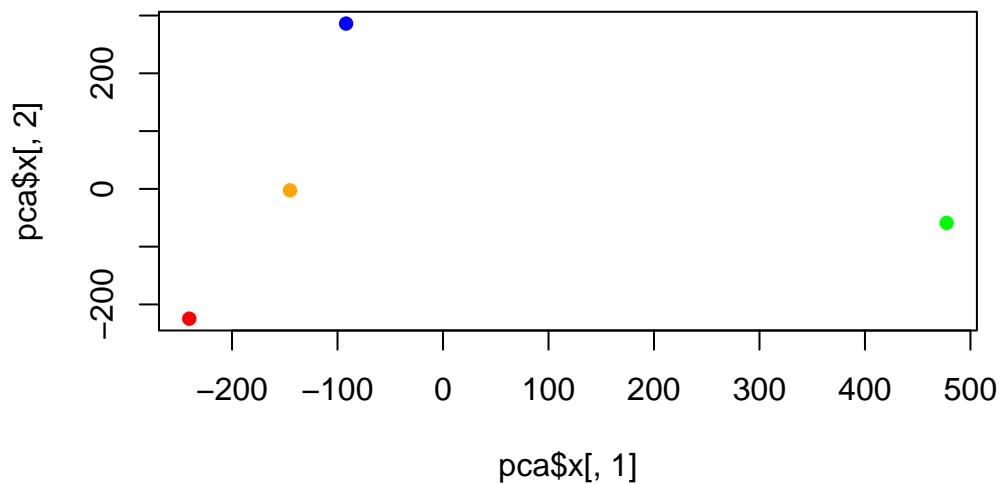
```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

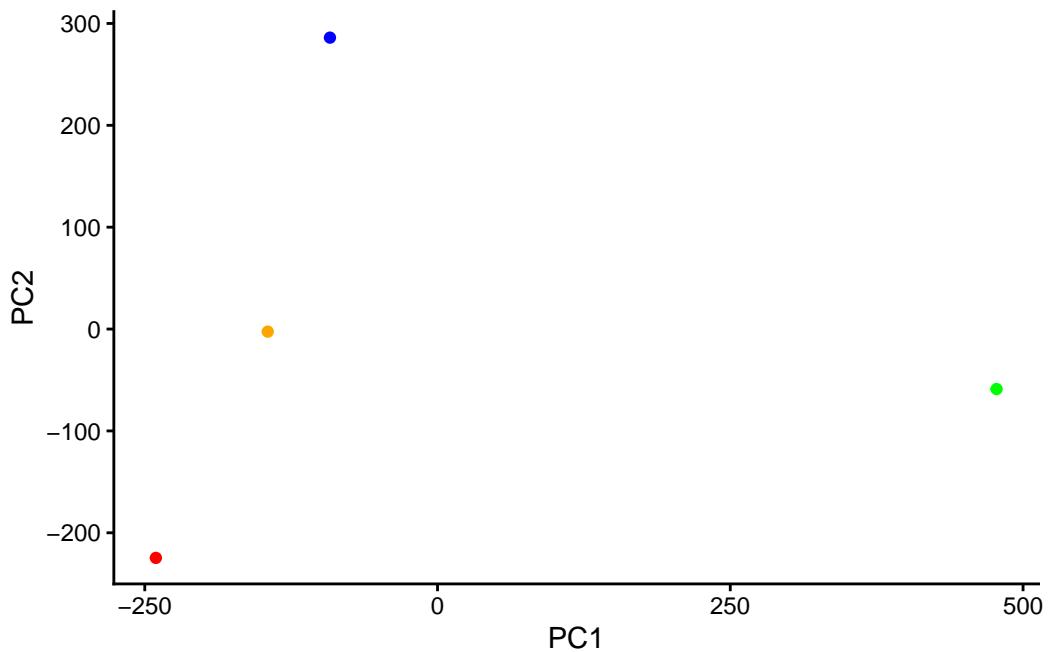
```
# See how PC1 and PC2 captured 96% of the variance -> good enough representative of the data
```

```
cols <- c("orange","red","blue","green")
plot(pca$x[,1], pca$x[,2], col=cols, pch=16)
```

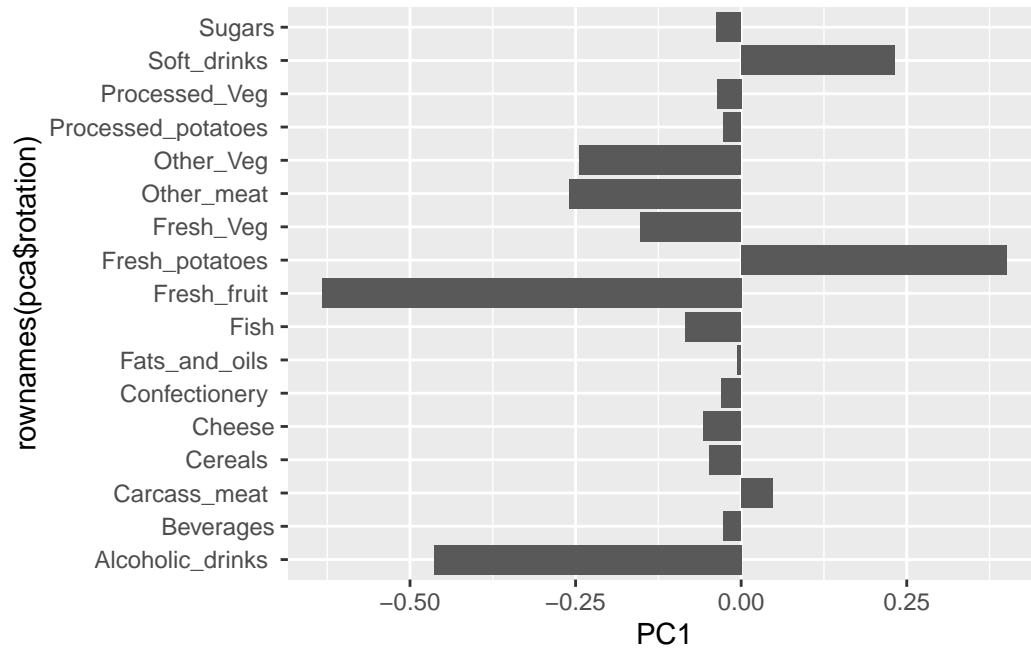


```
library(ggplot2)
```

```
ggplot(pca$x) +  
  aes(PC1, PC2) +  
  geom_point(col = cols) +  
  theme_classic()
```



```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```



```
# How much each variable weighs in the PCA
```

PCA looks really useful and we will come back to discuss this further next class (Class 8).